

Deploying a JHipster Microservice Cluster to Kubernetes on GCP

Abdelmajid Koubach
Mahmoud Mazouz

February 10, 2023

1 Introduction

Microservices have become increasingly popular in recent years as a way to build scalable and resilient applications. They allow for development teams to work on independent components, which can be deployed, updated, and managed separately. This approach results in a more flexible and efficient system, compared to monolithic architecture.

When deploying microservices, it is important to choose the right platform to support their needs. There are several cloud-based platforms available, such as Amazon Web Services (AWS), Google Cloud Platform (GCP), and Heroku. In this report, we will focus on deploying a JHipster microservice cluster to the Kubernetes engine in GCP.

2 JHipster and Microservices

JHipster is a popular open-source platform for generating, developing, and deploying modern web applications and microservices. It provides a pre-configured set of technologies, including Angular or React for the front-end, Spring Boot for the back-end, and either MySQL or MongoDB for the database.

In this project, we will deploy a JHipster microservice cluster, which consists of multiple independent microservices that communicate with each other to provide a complete application. Each microservice will be deployed as a separate container in a Kubernetes cluster, allowing for easy scaling and management.

3 Generation of the microservices architecture

The generation of each microservice with Jhipster went mostly without issues since they were launched independently using **docker-compose** and **gradlew**, and all services (Invoice, Notification, ProductOrder) were detected by the jhipster-registry service successfully. That's for the part without entities, as

for the refactoring of the application schema we chose the following patterns to define the **sku**, **upc**, **imageCdnUrl**, **thumbnailCdnUrl** of the ProductOrder entity:

Listing 1: JHipster Domain Language code

```
sku String required pattern (/^[a-zA-Z0-9]{1,50}$/)
upc String required pattern (/^[0-9]{12}$/)
...
/** For the content delivery network */
imageCdnUrl String pattern (/^https?:\/\/(?:[a-zA-Z0-9.-]*\.)?(?:cdn|cdn\.(?:[a-zA-Z]{2,})\.(?:com|net|org))\/.*\.(?:jpg|jpeg|png|gif|bmp|tiff)$/ )
...
/** For the content delivery network */
thumbnailCdnUrl String pattern (/^https?:\/\/.*\/thumb\/.*\.(?:jpg|jpeg|png|gif|bmp|tiff)$/ )
```

4 Microservices with Docker-compose

However for the docker-compose part where everything is bundled in a single group of containers the **notification** microservice fails to stay up because of an initialization issue with the MongoDB cluster mode since by default neither the replica sets or the config servers are initialized. So we had to manually execute and evaluate some initialization commands inside the containers and re-launch the notification container. The details for the initialization process are available on this Jhipster link : <https://www.jhipster.tech/docker-compose/#mongodb-cluster-mode>

One issue we could not solve in this part is monitoring. The services for Prometheus and Grafana are properly up, but they cannot communicate with the gateway server throwing the error **connection refused**. Moreover, unlike the part where every service is launched separately in the docker-compose bundle there are no default dashboards for the grafana service to use.

5 Microservices deployed on GCP

Deployment on the Google cloud platform requires a good **Kubernetes** cluster so that each machine has enough resources to sustain its pods and deployed microservice/s. We chose the following configuration :

Any other configurations were left on default. When the cluster is ready, all that is left to do is to push the images of the microservices onto the container registry and launch them as services on the cluster using the provided Jhipster shell program **./kubectl-apply.sh -f** (-f is the default option for the command as it requires an argument to execute).

Name of the cluster	micro-cluster
Auto-scalable cluster	YES
Number of nodes	5
Prometheus feature enabled	YES
Machine type	Standard (2 cpu, 8GB RAM)

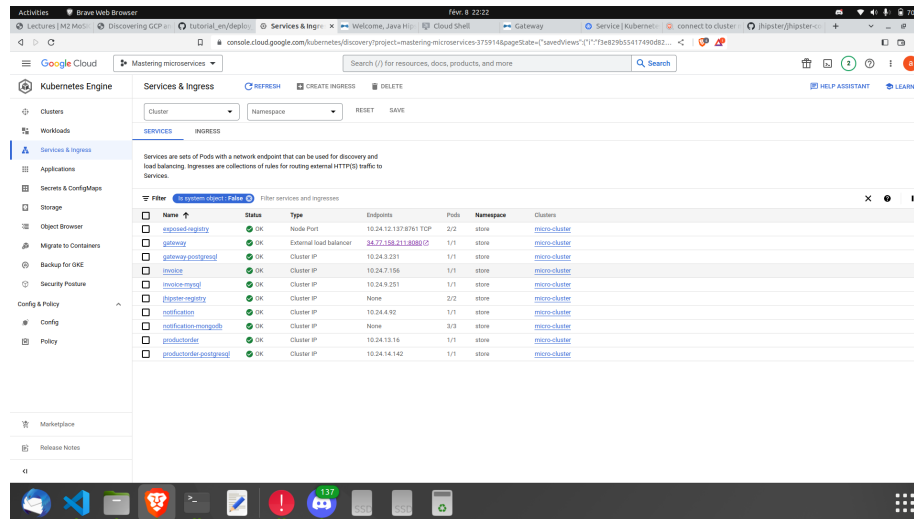


Figure 1: Deployment of microservices on GCP


The result looks like figure 1 and 2. We also enabled horizontal pod autoscaling for a service (Invoice) as illustrated in figure 3. The GCP dashboard while the services are deployed looked like figure 4.

6 Load Injection with Gatling

Instead of the offered version for Gatling (2.3.1) we used the latest version (3.9.0) because of a few reasons :

- An issue with **gatling.sh** which did not execute properly and throws errors
- The UI interface for **recorder.sh** offers the option of specifying the output type (java 8, java 11, scala, etc ...)
- The default outputs of version 2.3.1 is in scala and it produces some errors in the compilation phase

The result of the load injection with Gatling looks as shown in figure 5. The report is available as well as the simulation used in the Github repository (in the folder gatling-charts-highcharts-bundle-3.9.0-bundle). We didn't notice any



4

```

Events: <none>
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ # Enabling autoscaler for Invoice
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl autoscale replicaset invoice --cpu-percent=60 --min=2 --max=5
Error from server (NotFound): replicaset.apps "invoice" not found
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl autoscale replicaset invoice --cpu-percent=60 --min=2 --max=5 -n store
Error from server (NotFound): replicaset.apps "invoice" not found
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl autoscale ReplicaSet invoice --cpu-percent=60 --min=2 --max=5 -n store
Error from server (NotFound): replicaset.apps "invoice" not found
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl get rs
No resources found in default namespace.
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl get rs -n store
NAME                                DESIRED  CURRENT  READY  AGE
gateway-66646bd789                  1        1        1      68m
gateway-postgresql-6dc6655878       1        1        1      68m
invoice-7c97976788                  1        1        1      68m
invoice-mysql-646888ff55             1        1        1      68m
jhipster-prometheus-operator-6d9f9cddf 1        1        1      68m
notification-6466d9b988             1        1        1      68m
productorder-9d6f9ffff8             1        1        1      68m
productorder-postgresql-68d479f4d9   1        1        1      68m
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ # Enabling autoscaler for Invoice
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl autoscale rs invoice --cpu-percent=60 --min=2 --max=5 -n store
Error from server (NotFound): replicaset.apps "invoice" not found
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$ kubectl autoscale rs invoice-7c97976788 --cpu-percent=60 --min=2 --max=5 -n store
horizontalpodautoscaler.autoscaling/invoice-7c97976788 autoscaled
abdelmajid_koubach@cloudshell:~ (mastering-microservices-375914)$

```

Figure 3: Horizontal autoscaling for Invoice microservice

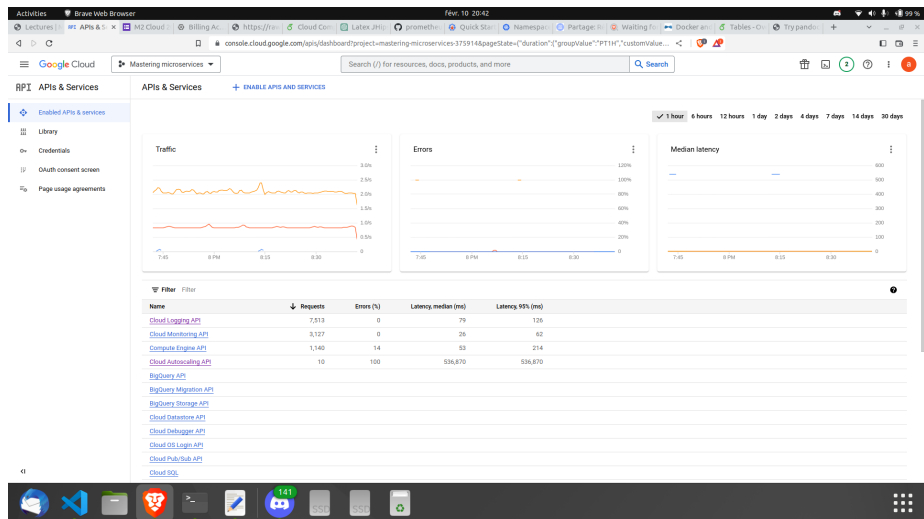


Figure 4: GCP dashboard

```

=====
---- Global Information -----
> request count                44 (OK=44    KO=0    )
> min response time            59 (OK=59    KO=-   )
> max response time            1239 (OK=1239 KO=-   )
> mean response time           207 (OK=207   KO=-   )
> std deviation                257 (OK=257   KO=-   )
> response time 50th percentile 110 (OK=110   KO=-   )
> response time 75th percentile 175 (OK=175   KO=-   )
> response time 95th percentile 716 (OK=716   KO=-   )
> response time 99th percentile 1211 (OK=1211 KO=-   )
> mean requests/sec            0.364 (OK=0.364 KO=-   )
---- Response Time Distribution -----
> t < 800 ms                   42 ( 95%)
> 800 ms <= t < 1200 ms        1 (  2%)
> t ≥ 1200 ms                   1 (  2%)
> failed                        0 (  0%)
=====

```

Figure 5: Running a Gatling recorder simulation on the gateway